

# Choosing the right optimizer


**Default: start with 'adam' (Adaptive Moment Estimation)**

- adam has advantages: adaptive 'lr' and momentum
- Use SGD when you need interpretability of optimization trajectory, have specific convergence goals, or want simplicity.
- The trends for SGD, Mini-Batch GD, and Full GD are well-characterized in terms of practical behavior (averaging vs. noise effects), but the underlying science is not fully understood.


# Let's visualize the training process

## Animation showing how NN learns predictions over epochs

NN Fitting ( $x^2 \bmod 5$ ), Epoch 283



- for animation, save NN's predictions at each training epoch.
- 'Callbacks' in Keras can do this for you. Add a chunk of code `model.fit()`, which Keras calls at certain points.



In what sense is the result (right panel) “correct”? Because it fits the training data - that’s all it knows.

Between points, NN interpolates smoothly, not because it understands but because it's the simplest way to reduce error.

Generalization isn't guaranteed. Outside training range, DNN behavior is unpredictable. Without more data, there's no true “right answer,” only guesses.


# Training | Validation | Test

## Who designed this building?

**Training:** learning via backprop & weight updates

**Validation:** Fine tune, prevent overfit; weights not updated

**Test:**  
real-world evaluation



Gehry =  
Shiny silver?

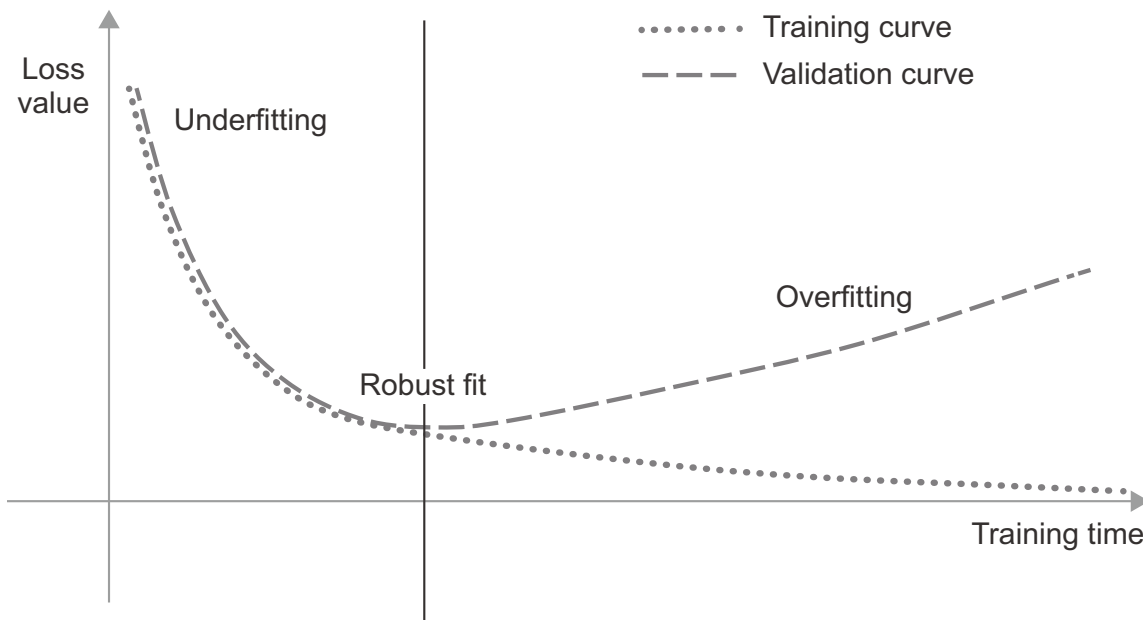


not all silver  
but curvy and  
irregular



# Underfit, robust fit & overfit

**"With four parameters I can fit an elephant, and with five I can make him wiggle his trunk." - Enrico Fermi**



**Underfitting:** The model is too simple to capture patterns —> high bias.

**Overfitting:** The model is too complex and memorizes training data —> poor generalization.

**Robust Fit:** A balanced model with optimal generalization with minimal bias and variance.

Quiz: You are training NN and validation loss starts increasing while the training loss keeps decreasing. What should you do?


- A. Increase # of epochs and train longer, hoping that it will generalize better.
- B. Reduce the model complexity or add regularization (e.g., dropout).
- C. Early stop: Go back to an earlier epoch where the validation loss was lowest and use early stopping to select that model.
- D. Reduce dataset size to prevent the model from learning too many unnecessary patterns.



# Why not just have training vs. test sets?

## Why do we need validation set?


Horse-picture from Pascal VOC data set



Source tag  
present



Classified  
as horse



No source  
tag present



Not classified  
as horse

C. Lothar Lenz  
www.pferdephotoarchiv.de

**Unmasking Clever Hans predictors and  
assessing what machines really learn**

<https://www.nature.com/articles/s41467-019-08987-4>

- Hyperparameter tuning (# of neurons, layers, activation func) relies on validation sets.
- Over time, DNN may unknowingly become artificially overfitted to the validation data.
- Example (left): A model trained on Pascal VOC dataset learned to recognize watermarks rather than horses. Well, the model wasn't completely wrong since it found a correlation: "pferde" (horses in German) often appeared in horse images. But this is shortcut learning, not generalization.
- Lesson: We must test on never-seen new data that is separate to ensure true generalization.


# Dropout: Inspired by biology

```
layer = nn.Sequential(nn.Linear(512, 256), nn.ReLU(), nn.Dropout(p=0.2))
```

Dropout: A Simple Way to Prevent Neural Networks from Overfitting

Nitish Srivastava  
Geoffrey Hinton  
Alex Krizhevsky  
Ilya Sutskever  
Ruslan Salakhutdinov

NITISH@CS.TORONTO.EDU  
HINTON@CS.TORONTO.EDU  
KRIZ@CS.TORONTO.EDU  
ILYA@CS.TORONTO.EDU  
RSALAKHU@CS.TORONTO.EDU

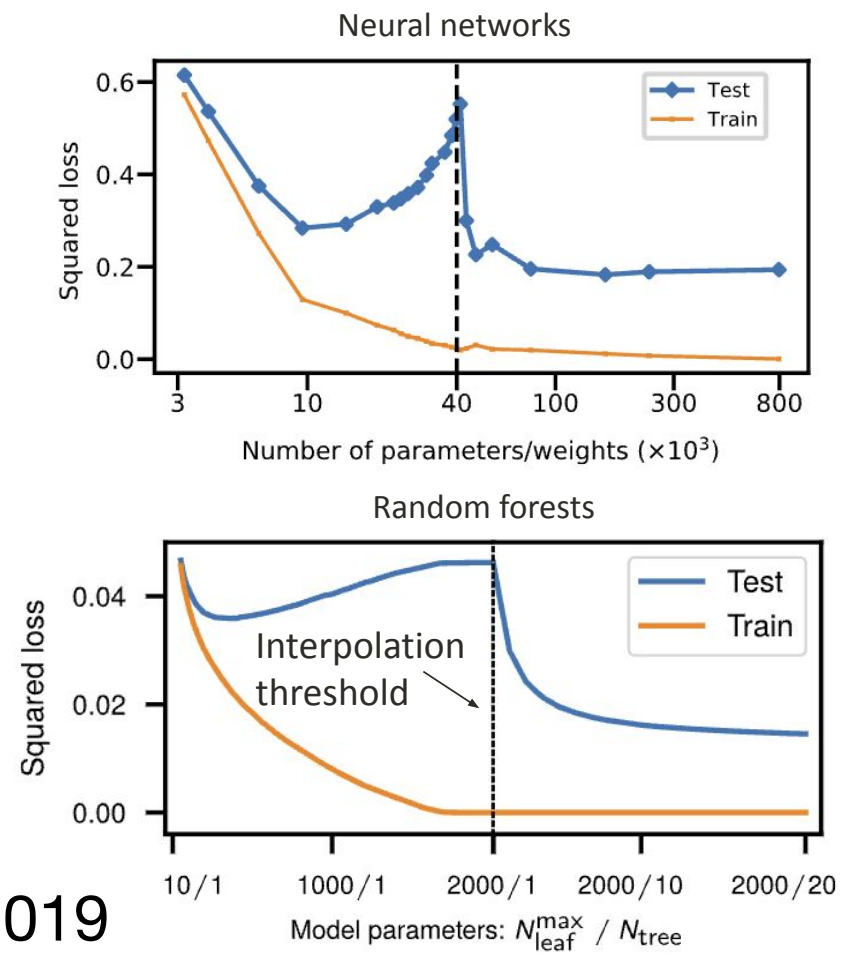


“A motivation for dropout comes from a theory of the role of sex in evolution.”

- Hinton’s sexual reproduction analogy: Just as genetic crossover prevents the inheritance of harmful mutations, dropout randomly removes neurons during training, forcing NNs to develop redundant, robust features since it can’t rely on specific neurons being active.
- Implicit averaging: A network with dropout can be seen as an ensemble of many sub-networks, improving stability and predictive accuracy.
- Practical tip: Use 20–50% dropout in fully connected layers; convolutional layers require less due to their intrinsic weight sharing.


# Beyond Overfitting: Double Descent

- Interpolation threshold & Generalization: When model capacity exceeds a critical point, test error initially increases due to memorization but later decreases, improving generalization beyond overfitting regime.



## DEEP DOUBLE DESCENT: WHERE BIGGER MODELS AND MORE DATA HURT

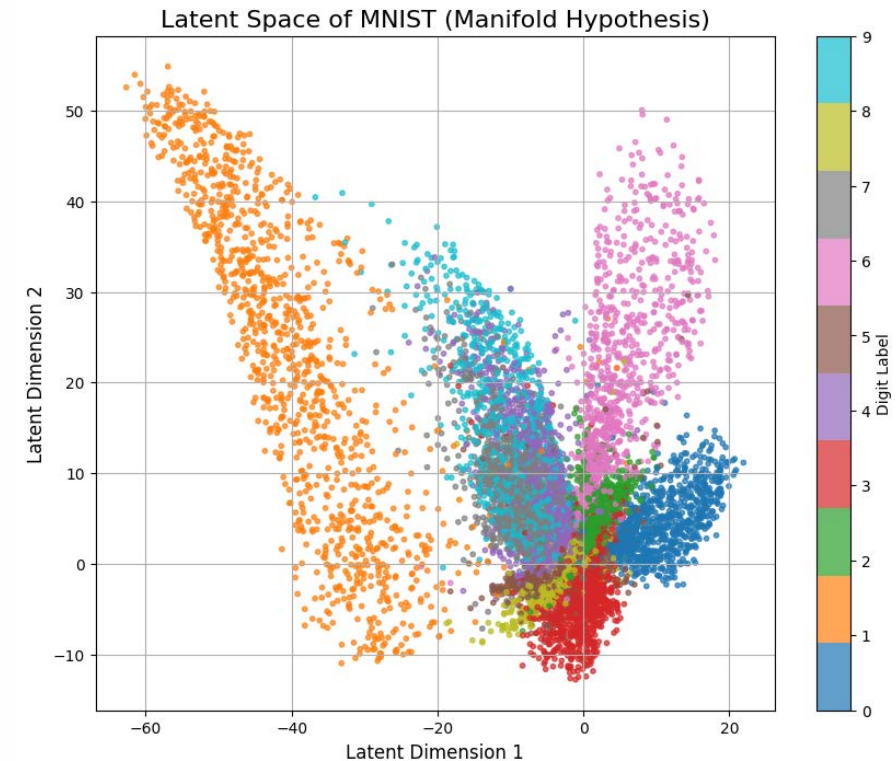


<b>Preetum Nakkiran*</b> Harvard University	<b>Gal Kaplun†</b> Harvard University	<b>Yamini Bansal†</b> Harvard University	<b>Tristan Yang</b> Harvard University
<b>Boaz Barak</b> Harvard University	<b>Ilya Sutskever</b> OpenAI		



# Manifold Hypothesis again: Why does deep learning work?

## Why does generalization work? Is it just interpolation?

- High-dimensional data, like images and text, often lie on smooth and differentiable manifolds, making them amenable to deep learning. Neural nets exploit this structure to reduce complexity, capturing meaningful variations rather than memorizing raw data. This shows up in Variational Autoencoders (VAEs), which map inputs onto a continuous, low-dimensional latent space, enabling smooth transformations.





# Can deep learning discover sorting rules?

✓  
4m



```
import numpy as np
from tensorflow import keras

X = np.random.randint(0, 100, (50000, 10)) # Training sample
Y = np.sort(X, axis=1) # Sorted input

model = keras.Sequential([
    keras.layers.Dense(32, activation='relu', input_shape=(10,)),
    keras.layers.Dense(64, activation='relu'),
    keras.layers.Dense(32, activation='relu'),
    keras.layers.Dense(10) # Output layer
])

model.compile(optimizer='adam', loss='mse')
model.fit(X, Y, epochs=100, batch_size=64)

test_samples = np.array([[3, 15, 9, 1, 20, 32, 13, 23, 55, 21],
                          [99, 10, 75, 2, 2, 52, 89, 10, 32, 2],
                          [0, 99, 50, 50, 1, 77, 33, 88, 22, 33]])

predictions = model.predict(test_samples).round().astype(int)
print("Predicted:\n", predictions)
print("Sorted:\n", np.sort(test_samples, axis=1))
```

# General strategy for DNN architecture design

## How many hidden layers and neurons? Which activation func?

1. Layer structure: arrange neurons in hidden layers to graduate increase (to capture more complexity for feature extraction) or decrease (toward compression, decision making).
  2. Width (neurons per layer): Start small and increase # of neurons as needed. Excessive neurons => overfit risk.
  3. Depth (# of hidden layers): Start with 2-4 layers and scale as needed. Deeper architecture allows more abstract representation (good) but training is more difficult.
  4. Activation function: Use ReLU by default to avoid saturation. Tanh or sigmoid are useful in probabilities or bounded outputs, but suffer from vanishing gradients in deep nets.
  5. Dropout and regularization: Add dropout (20—50% rate) to prevent overfit;
- Occam's razor: the simplest DNN architecture giving the desired result is the best.

# End-to-end MNIST

```
import tensorflow as tf

(x_train, y_train), (x_test, y_test) = tf.keras.datasets.mnist.load_data()
x_train, x_test = x_train / 255.0, x_test / 255.0 # Why normalize pixel values?

model = tf.keras.Sequential([
    tf.keras.layers.Flatten(input_shape=(28, 28)), # Input layer
    tf.keras.layers.Dense(32, activation='relu'), # Hidden layer
    tf.keras.layers.Dropout(0.2), # what's the purpose of this?
    tf.keras.layers.Dense(10, activation='softmax') # why is softmax used?
])

model.compile(optimizer=tf.keras.optimizers.SGD(learning_rate=0.01),
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])
# change batch_size and see how it affects training speed and accuracy
history = model.fit(x_train, y_train, batch_size=128, epochs=20, verbose=1)
test_loss, test_acc = model.evaluate(x_test, y_test, verbose=2)
```